

Nonsmooth Optimization on Riemannian Manifolds

Ronny Bergmann

NTNU, Trondheim, Norway.

Adaptive Bayesian Intelligence Seminar

RIKEN AIP, Tokyo, Japan, February 18, 2026.



Nonsmooth Optimization on Riemannian Manifolds

NTNU

We are looking for **numerical algorithms** to find

$$\arg \min_{p \in \mathcal{M}} f(p)$$

where

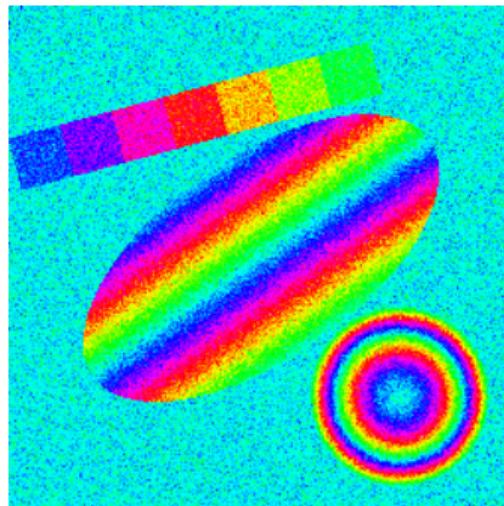
- ▶ \mathcal{M} is a Riemannian manifold
- ▶ $f: \mathcal{M} \rightarrow \bar{\mathbb{R}}$ is a function
- ⚠ f might be **nonsmooth** and/or **nonconvex**
- ⚠ \mathcal{M} might be **high-dimensional**
- 💡 f has some “nice structure”

Manifold-valued signal and image processing

- ▶ variational models for
denoising, inpainting, deconvolution, segmentation, ...
- ▶ applications in medical imaging, computer vision
- ▲ nonlinear (non-Euclidean) data

Examples

- ▶ phase-valued data (\mathbb{S}^1)
- ▶ wind-fields, GPS (\mathbb{S}^2)
- ▶ DT-MRI ($\mathcal{P}(3)$)
- ▶ EBSD, (grain) orientations ($\text{SO}(n)$)



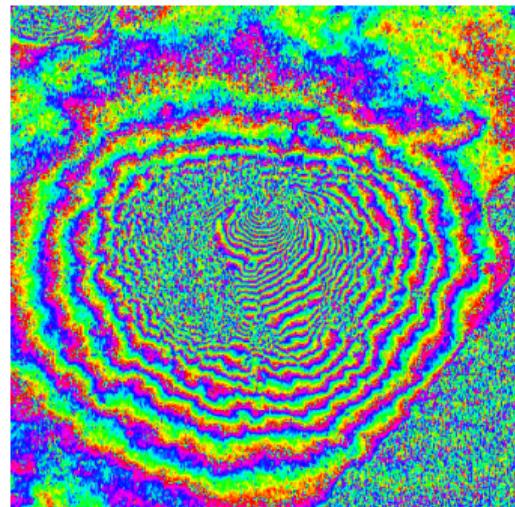
Artificial noisy phase-valued data.

Manifold-valued signal and image processing

- ▶ variational models for
denoising, inpainting, deconvolution, segmentation, ...
- ▶ applications in medical imaging, computer vision
- ▲ nonlinear (non-Euclidean) data

Examples

- ▶ phase-valued data (\mathbb{S}^1)
- ▶ wind-fields, GPS (\mathbb{S}^2)
- ▶ DT-MRI ($\mathcal{P}(3)$)
- ▶ EBSD, (grain) orientations ($\text{SO}(n)$)



InSAR-Data of Mt. Vesuvius.

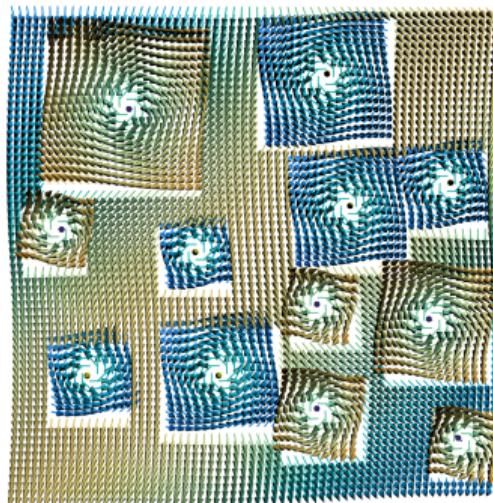
[Rocca, Prati, Guarnieri, 1997]

Manifold-valued signal and image processing

- ▶ variational models for
denoising, inpainting, deconvolution, segmentation, ...
- ▶ applications in medical imaging, computer vision
- ▲ nonlinear (non-Euclidean) data

Examples

- ▶ phase-valued data (\mathbb{S}^1)
- ▶ wind-fields, GPS (\mathbb{S}^2)
- ▶ DT-MRI ($\mathcal{P}(3)$)
- ▶ EBSD, (grain) orientations ($\text{SO}(n)$)



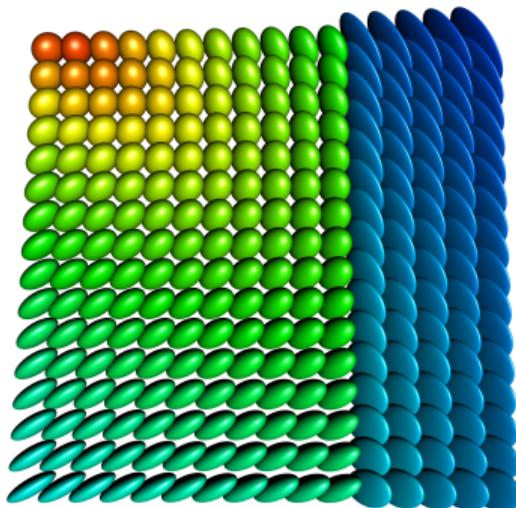
Artificial noisy data on the sphere \mathbb{S}^2 .

Manifold-valued signal and image processing

- ▶ variational models for
denoising, inpainting, deconvolution, segmentation, ...
- ▶ applications in medical imaging, computer vision
- ▲ nonlinear (non-Euclidean) data

Examples

- ▶ phase-valued data (\mathbb{S}^1)
- ▶ wind-fields, GPS (\mathbb{S}^2)
- ▶ DT-MRI ($\mathcal{P}(3)$)
- ▶ EBSD, (grain) orientations ($\text{SO}(n)$)



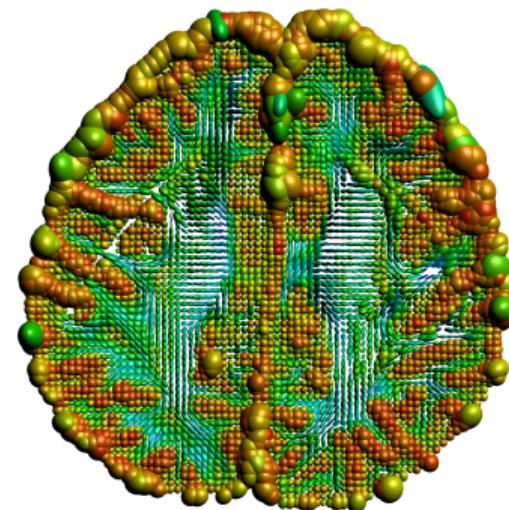
Artificial diffusion data,
each pixel is a sym. pos. def. matrix.

Manifold-valued signal and image processing

- ▶ variational models for
denoising, inpainting, deconvolution, segmentation, ...
- ▶ applications in medical imaging, computer vision
- ▲ nonlinear (non-Euclidean) data

Examples

- ▶ phase-valued data (\mathbb{S}^1)
- ▶ wind-fields, GPS (\mathbb{S}^2)
- ▶ DT-MRI ($\mathcal{P}(3)$)
- ▶ EBSD, (grain) orientations ($\text{SO}(n)$)



DT-MRI of the human brain.

Camino Project: cmic.cs.ucl.ac.uk/camino

Manifold-valued signal and image processing

- ▶ variational models for
denoising, inpainting, deconvolution, segmentation, ...
- ▶ applications in medical imaging, computer vision
- ▲ nonlinear (non-Euclidean) data

Examples

- ▶ phase-valued data (\mathbb{S}^1)
- ▶ wind-fields, GPS (\mathbb{S}^2)
- ▶ DT-MRI ($\mathcal{P}(3)$)
- ▶ EBSD, (grain) orientations ($\text{SO}(n)$)



Grain orientations in EBSD data.

MTEX toolbox: mtex-toolbox.github.io



Constraints and/or geometry

NTNU

constraints

- ▶ needs an embedding
- ▶ might not always yield a manifold
- 😊 slightly more flexible
- 🙁 algorithms have to deal with constraints
- 🙁 results might be infeasible

geometry

- ▶ might work agnostic of an embedding
- 😊 quotient manifolds
- 😊 we can use any unconstrained algorithm...
- 🙁 ...after adapting it to the manifold setting
- 😊 algorithms stay on the manifold
- ➡ always feasible

We can also consider a combination of both:
constrained optimization on manifolds.

[Liu, Boumal, 2019; RB, Herzog, 2019]

A Riemannian Manifold \mathcal{M}

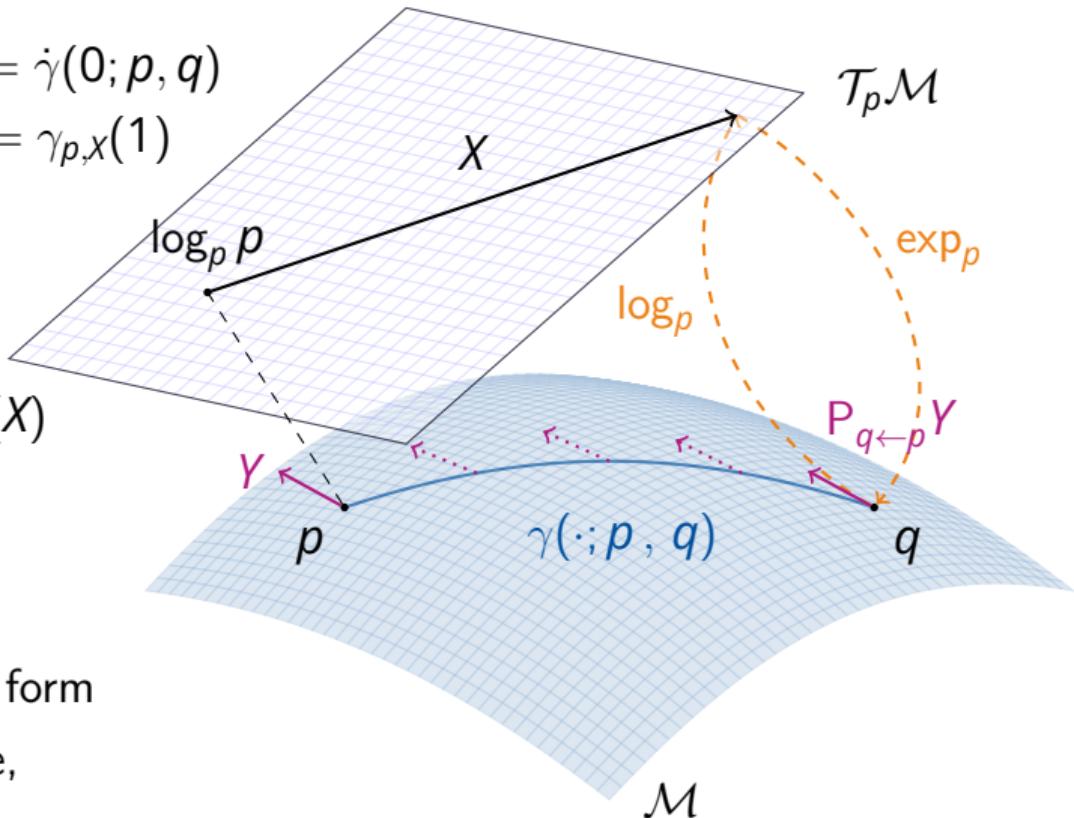
Notation.

- ▶ Logarithmic map $\log_p q = \dot{\gamma}(0; p, q)$
- ▶ Exponential map $\exp_p X = \gamma_{p,X}(1)$
- ▶ Geodesic $\gamma(\cdot; p, q)$
- ▶ Tangent space $T_p\mathcal{M}$
- ▶ inner product $(\cdot, \cdot)_p$
- ▶ parallel transport $\text{PT}_{p \leftarrow q}(X)$
- ▶ distance function $d(p, q)$

Numerics.

$\exp_p, \log_p, \text{PT}_{p \leftarrow q}$ maybe
not avail, efficiently/in closed form

⇒ use a retraction, its inverse,
a vector transport instead



(Geodesic) Convexity



NTNU

[Sakai, 1996; Udriște, 1994]

A set $\mathcal{C} \subset \mathcal{M}$ is called (strongly geodesically) **convex**
if for all $p, q \in \mathcal{C}$ the geodesic $\gamma(\cdot; p, q)$ is unique and lies in \mathcal{C} .

A function $f: \mathcal{C} \rightarrow \overline{\mathbb{R}}$ is called (geodesically) **convex**
if for all $p, q \in \mathcal{C}$ the composition $f(\gamma(t; p, q)), t \in [0, 1]$, is convex.



The Riemannian Subdifferential

NTNU

Let \mathcal{C} be a convex set.

The **subdifferential** of f at $p \in \mathcal{C}$ is given by

[Ferreira, Oliveira, 2002; Lee, 2003; Udriște, 1994]

$$\partial_{\mathcal{M}} f(p) := \left\{ \xi \in \mathcal{T}_p^* \mathcal{M} \mid f(q) \geq f(p) + \langle \xi, \log_p q \rangle_p \text{ for } q \in \mathcal{C} \right\},$$

where

- ▶ $\mathcal{T}_p^* \mathcal{M}$ is the dual space of $\mathcal{T}_p \mathcal{M}$, also called **cotangent space**
- ▶ $\langle \cdot, \cdot \rangle_p$ denotes the duality pairing on $\mathcal{T}_p^* \mathcal{M} \times \mathcal{T}_p \mathcal{M}$
- ▶ numerically we use musical isomorphisms $X = \xi^\flat \in \mathcal{T}_p \mathcal{M}$ to obtain a subset of $\mathcal{T}_p \mathcal{M}$



The Proximal Point Algorithm

Euclidean case. For $f: \mathbb{R}^n \rightarrow \overline{\mathbb{R}}$, $\lambda > 0$, the proximal map given by

[Moreau, 1965; Rockafellar, 1970]

$$\text{prox}_{\lambda f}(x) = \arg \min_{y \in \mathbb{R}^n} \left\{ f(y) + \frac{1}{2\lambda} \|y - x\|^2 \right\}.$$

Riemannian case. For $f: \mathcal{M} \rightarrow \overline{\mathbb{R}}$, $\lambda > 0$, the proximal map is given by

[Ferreira, Oliveira, 2002]

$$\text{prox}_{\lambda f}(p) = \arg \min_{q \in \mathcal{M}} \left\{ f(q) + \frac{1}{2\lambda} d(p, q)^2 \right\}.$$

For both. A minimizer p^* of f is a fixed point for $\text{prox}_{\lambda f}$.

Proximal Point Algorithm (PPA). Given $p^{(0)} \in \mathcal{M}$, $\lambda_k > 0$, iterate

$$p^{(k+1)} = \text{prox}_{\lambda_k f}(p^{(k)}).$$



The Cyclic Proximal Point Algorithm

[Bertsekas, 2011; Bačák, 2014]

For a splitting $f(p) = \sum_{i=1}^c f_i(p)$ and some $p_0 \in \mathcal{M}$, we can use

$$p_{k+\frac{i+1}{c}} = \text{prox}_{\lambda_k f_i}(p_{k+\frac{i}{c}}), \quad i = 0, \dots, c-1, \quad k = 0, 1, \dots$$

On a Hadamard manifold \mathcal{M} : Convergence to a minimizer of f if

- ▶ all f_i proper, convex, lower semi-continuous
- ▶ $\{\lambda_k\}_{k \in \mathbb{N}} \in \ell_2(\mathbb{N}) \setminus \ell_1(\mathbb{N})$.
- ▶ also for
 - ▶ random order of the $\text{prox}_{\lambda f_i}$
 - ▶ inexact evaluations of the $\text{prox}_{\lambda f_i}$

[Bačák, RB, Steidl, Weinmann, 2016]

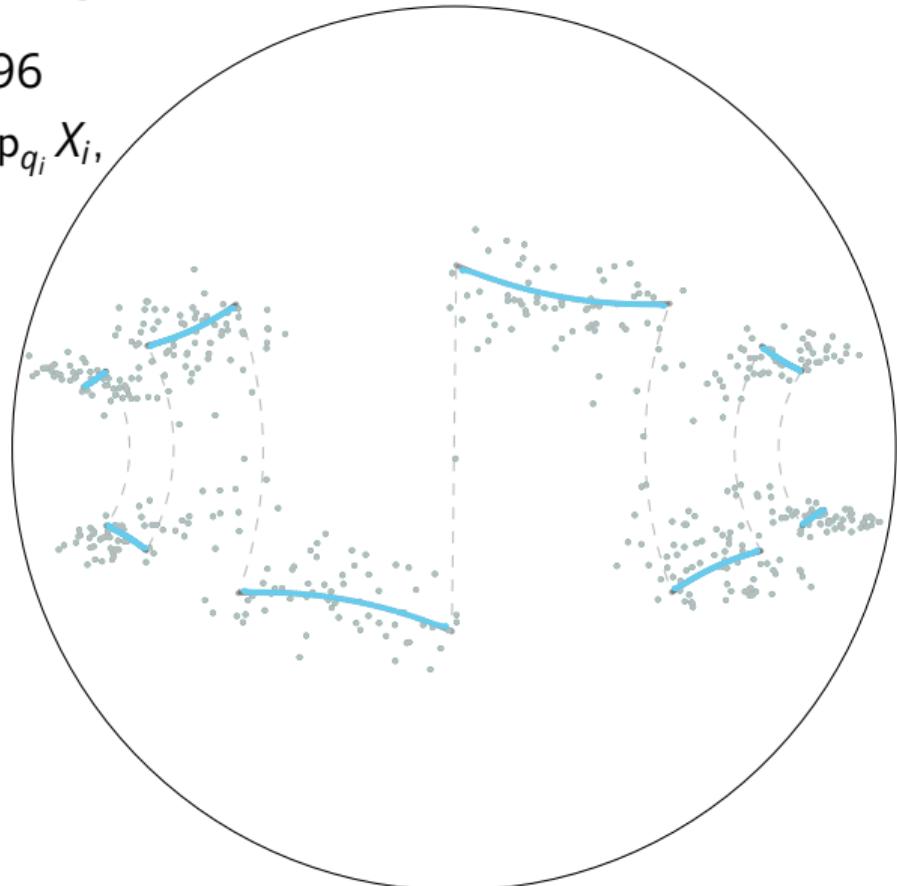
! no convergence rate

Denoising a Signal on Hyperbolic Space \mathcal{H}^2

- ▶ signal $q \in \mathcal{M}, (\mathcal{H}^2)^n, n = 496$
- ▶ noisy signal $\bar{q} \in \mathcal{M}, \bar{q}_i = \exp_{q_i} X_i, \sigma = 0.1$
- ▶ ROF Model:

$$\begin{aligned} \arg \min_{p \in \mathcal{M}} \frac{1}{n} d_{\mathcal{M}}(p, q)^2 \\ + \alpha \sum_{i=1}^{n-1} d_{\mathcal{H}^2}(p_i, p_{i+1}) \end{aligned}$$

- ▶ Setting $\alpha = 0.05$ yields reconstruction p^* .





Algorithms for Denoising a Signal

NTNU

- ▶ Riemannian Convex Bundle Method (RCBM) [RB, Herzog, Jasa, 2024]
- ▶ Proximal Bundle Algorithm (PBA) [Hoseini Monjezi, Nobakhtian, Pouryayevali, 2021]
- ▶ Subgradient Method (SGM) [Ferreira, Oliveira, 1998]
- ▶ Cyclic Proximal Point Algorithm (CPPA) [Bačák, 2014]

Algorithm	Iter.	Time (sec.)	Objective	Error
RCBM	3417	51.393	1.7929×10^{-3}	3.3194×10^{-4}
PBA	15 000	102.387	1.8153×10^{-3}	4.3874×10^{-4}
SGM	15 000	99.604	1.7920×10^{-3}	3.3080×10^{-4}
CPPA	15 000	94.200	1.7928×10^{-3}	3.3230×10^{-4}



The Douglas Rachford Algorithm

For a splitting $f = g + h$, where both are possibly nonsmooth, use the reflection at the proximal map

$$R_{\lambda f}(p) = \exp_{\text{prox}_{\lambda f}(p)}(-\log_{\text{prox}_{\lambda f}(p)}(p)) \quad (\text{Euclidean: } 2\text{ prox}_{\lambda f}(x) - x)$$

The Douglas Rachford algorithm reads for some $r^{(0)} \in \mathcal{M}$, $\eta > 0$

$$p^{(k)} = R_{\eta g}(r^{(k)})$$

[RB, Persch, Steidl, 2016]

$$q^{(k)} = R_{\eta h}(p^{(k)})$$

$$r^{(k+1)} = \gamma(\lambda_k; r^{(k)}, q^{(k)}) \quad (\gamma \text{ is a geodesic})$$

- ▶ converges on Hadamard manifolds if
 - ▶ g, h proper, convex, lsc.
 - ▶ $\lambda_k \in [0, 1]$ and $\sum_k \lambda_k(1 - \lambda_k) = \infty$
- ▶ ...to a fixed point of $R_{\lambda g} \circ R_{\lambda h}$ (in $r^{(k)}$)
- ▶ ...to a minimizer of f in the “shadow iterates” $\text{prox}_{\eta g}(r^{(k)})$



The Fenchel Conjugate

NTNU

The Fenchel conjugate of a function $f: \mathbb{R}^n \rightarrow \overline{\mathbb{R}}$ is given by

$$f^*(\xi) := \sup_{x \in \mathbb{R}^n} \langle \xi, x \rangle - f(x) = \sup_{x \in \mathbb{R}^n} \begin{pmatrix} \xi \\ -1 \end{pmatrix}^\top \begin{pmatrix} x \\ f(x) \end{pmatrix}$$

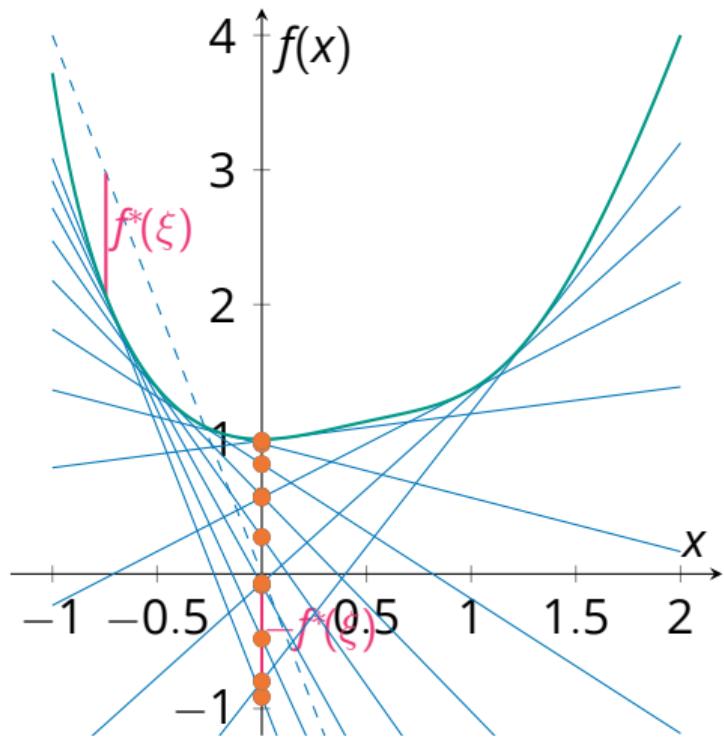
- ▶ given $\xi \in \mathbb{R}^n$: maximize the distance between $\xi^\top \cdot$ and f
- ▶ can also be written in the epigraph

The Fenchel biconjugate reads

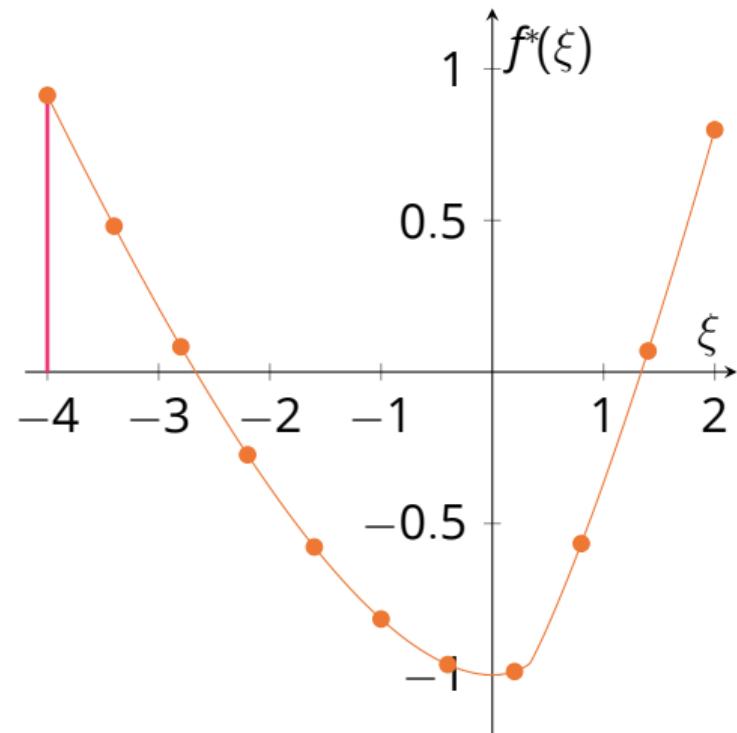
$$f^{**}(x) = (f^*)^*(x) = \sup_{\xi \in \mathbb{R}^n} \langle \xi, x \rangle - f^*(\xi).$$

Illustration of the Fenchel Conjugate

The function f



The Fenchel conjugate f^*





The (Riemannian) m -Fenchel Conjugate

[RB, Herzog, Silva Louzeiro, Tenbrinck, Vidal-Núñez, 2021]

Idea. Localize to $\mathcal{C} \subset \mathcal{M}$ around a point m which “acts as” 0.

The m -Fenchel conjugate of a function $f: \mathcal{C} \rightarrow \overline{\mathbb{R}}$ is given by

$$f_m^*(\xi_m) := \sup_{X \in \mathcal{L}_{\mathcal{C},m}} \{ \langle \xi_m, X \rangle - f(\exp_m X) \},$$

where $\mathcal{L}_{\mathcal{C},m} := \{X \in T_m \mathcal{M} \mid q = \exp_m X \in \mathcal{C} \text{ and } \|X\|_p = d(q, p)\}$.

Let $m' \in \mathcal{C}$. The mm' -Fenchel-biconjugate $F_{mm'}^{**}: \mathcal{C} \rightarrow \overline{\mathbb{R}}$ is given by

$$F_{mm'}^{**}(p) = \sup_{\xi_{m'} \in T_{m'}^* \mathcal{M}} \{ \langle \xi_{m'}, \log_{m'} p \rangle - F_m^*(P_{m \leftarrow m'} \xi_{m'}) \},$$

where usually we only use the case $m = m'$.

The exact Riemannian Chambolle–Pock Algorithm

[RB, Herzog, Silva Louzeiro, Tenbrinck, Vidal-Núñez, 2021; Valkonen, 2014; Chambolle, Pock, 2011]

To solve

$$\arg \min_{p \in \mathcal{C} \subset \mathcal{M}} \{f(p) + g(\Lambda(p))\},$$

Input: $m, p^{(0)} \in \mathcal{C} \subset \mathcal{M}$, $n = \Lambda(m)$, $\xi_n^{(0)} \in \mathcal{T}_n^*\mathcal{N}$, and $\sigma, \tau, \theta > 0$

1: $k \leftarrow 0$

2: $\bar{p}^{(0)} \leftarrow p^{(0)}$

3: **while** not converged **do**

4: $\xi_n^{(k+1)} \leftarrow \text{prox}_{\tau g_n^*}(\xi_n^{(k)} + \tau (\log_n \Lambda(\bar{p}^{(k)}))^{\flat})$

5: $p^{(k+1)} \leftarrow \text{prox}_{\sigma f}\left(\exp_{p^{(k)}}\left(P_{p^{(k)} \leftarrow m}(-\sigma D\Lambda(m)^*[\xi_n^{(k+1)}])^\sharp\right)\right)$

6: $\bar{p}^{(k+1)} \leftarrow \exp_{p^{(k+1)}}(-\theta \log_{p^{(k+1)}} p^{(k)})$

7: $k \leftarrow k + 1$

8: **end while**

Output: $p^{(k)}$



Proximal Gradient

NTNU

For a splitting $f = g + h$, where g is smooth and h is possibly nonsmooth, both are convex.

The proximal gradient method reads for given $p^{(0)} \in \mathcal{M}$, $\lambda_k \in (0, \frac{1}{L}]$ reads

[RB, Jasa, John, Pfeffer, 2025b]

$$p^{(k+1)} = \text{prox}_{\lambda_k h}(\exp_{p^{(k)}}(-\lambda_k \text{grad } g(p^{(k)}))).$$

- ▶ convergence rates: sublinear (convex) linear (strongly convex)
 - ▶ a generalization of the prox-grad inequality
 - ▶ even the **nonconvex** case: sublinear convergence to ε -stationary points
- [RB, Jasa, John, Pfeffer, 2025a]
- ! though here: proximal map maybe not unique minimizer



NTNU

The Riemannian DC Algorithm

To solve a Difference of Convex problem

[RB, Ferreira, Santos, Souza, 2024]

$$\arg \min_{p \in \mathcal{M}} g(p) - h(p).$$

use

The Riemannian Difference of Convex Algorithm.

Input: An initial point $p^{(0)} \in \text{dom}(g)$, g and $\partial_{\mathcal{M}} h$

1: Set $k = 0$.

2: **while** not converged **do**

3: Take $X^{(k)} \in \partial_{\mathcal{M}} h(p^{(k)})$

4: Compute the next iterate $p^{(k+1)}$ as

$$p^{(k+1)} \in \arg \min_{p \in \mathcal{M}} g(p) - (X^{(k)}, \log_{p^{(k)}} p)_{p^{(k)}}.$$

5: Set $k \leftarrow k + 1$

6: **end while**

Convergence of the Riemannian DCA

Let $\{p^{(k)}\}_{k \in \mathbb{N}}$ and $\{X^{(k)}\}_{k \in \mathbb{N}}$ be the iterates and subgradients of the RDCA.

Theorem.

[RB, Ferreira, Santos, Souza, 2024]

If \bar{p} is a cluster point of $\{p^{(k)}\}_{k \in \mathbb{N}}$, then $\bar{p} \in \text{dom}(g)$ and there exists a cluster point \bar{X} of $\{X^{(k)}\}_{k \in \mathbb{N}}$ s.t. $\bar{X} \in \partial g(\bar{p}) \cap \partial h(\bar{p})$.

⇒ Every cluster point of $\{p^{(k)}\}_{k \in \mathbb{N}}$, if any, is a critical point of f .

Proposition.

[RB, Ferreira, Santos, Souza, 2024]

Let g be σ -strongly (geodesically) convex. Then

$$f(p^{(k+1)}) \leq f(p^{(k)}) - \frac{\sigma}{2} d^2(p^{(k)}, p^{(k+1)})$$

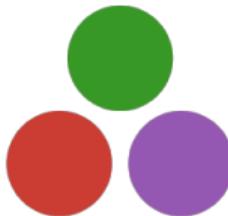
and $\sum_{k=0}^{\infty} d^2(p^{(k)}, p^{(k+1)}) < \infty$, so in particular $\lim_{k \rightarrow \infty} d(p^{(k)}, p^{(k+1)}) = 0$.



NTNU

Software

Goals of the Software – Why Julia?



Goals.

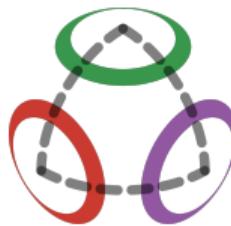
- ▶ abstract definition of manifolds and optimization thereon
- ⇒ implement abstract solvers on a generic manifold
- ▶ well-documented and well-tested
- ▶ fast.
- ⇒ “Run your favourite solver on your favourite manifold”.

Why ⚙️ Julia?

julialang.org

- ▶ high-level language, properly typed
- ▶ multiple dispatch, e. g. `*(::AbstractMatrix, ::AbstractMatrix)`
- ▶ just-in-time compilation, solves two-language problem
 - ⇒ “nice to write” and as fast as C/C++
- ▶ I like the community

ManifoldsBase.jl – Motivation



Goal. Provide a generic interface to manifolds for

- ▶ defining own (new) manifolds
- ▶ implementing **generic** algorithms on an arbitrary manifold \mathcal{M}

A Manifold. a Riemannian manifold is a subtype of `AbstractManifold{F}`

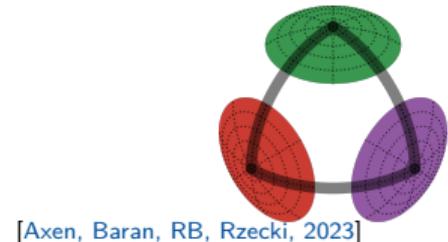
- ▶ $F \in \{\mathbb{R}, \mathbb{C}, \mathbb{H}\}$: field the manifold is build on
- ▶ stores all “general” information, (mainly) the manifold dimension
- ▶ example (from `Manifolds.jl`): `M = Sphere(2)`

Points and Tangent vectors.

- ▶ by default not typed, often `<:AbstractArray`
- ▶ we provide `<:AbstractManifoldPoint` and `<:TVector` for more advanced ones

Manifolds.jl

Goal. Provide a library of Riemannian manifolds, that is efficiently implemented and well-documented



[Axen, Baran, RB, Rzecki, 2023]

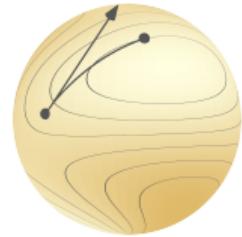
Meta. generic implementations for $\mathcal{M}^{n \times m}$, $\mathcal{M}_1 \times \mathcal{M}_2$, vector- and tangent-bundles, esp. $T_p\mathcal{M}$, or Lie groups

Library. Implemented functions for

- ▶ Circle, Sphere, Torus, Hyperbolic, Projective Spaces, Hamiltonian
- ▶ (generalized, symplectic) Stiefel, Rotations
- ▶ (generalized, symplectic) Grassmann, fixed rank matrices
- ▶ Symmetric Positive Definite matrices, with fixed determinant
- ▶ (several) Multinomial, (skew-)symmetric, and symplectic matrices
- ▶ Tucker & Oblique manifold, Kendall's Shape space
- ▶ probability simplex, orthogonal and unitary matrices, ...

Manopt.jl

Goal. Provide optimization algorithms on Riemannian manifolds.



Features. Given a `Problem p` and a `SolverState s`,
implement `initialize_solver!(p, s)` and `step_solver!(p, s, i)`
⇒ an algorithm in the `Manopt.jl` interface

Highlevel interfaces like `gradient_descent(M, f, grad_f)`
on any manifold `M` from `Manifolds.jl`.

All provide `debug` output, recording, cache & `counting` capabilities,
as well as a library of step sizes and `stopping criteria`.

Manopt family.



manoptjl.org

[RB, 2022]



manopt.org

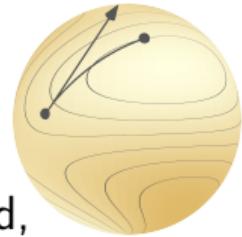
[Boumal, Mishra, Absil, Sepulchre, 2014]



pymanopt.org

[Townsend, Koep, Weichwald, 2016]

List of Algorithms in Manopt.jl



Derivative-Free Nelder-Mead, Particle Swarm, CMA-ES, MADS

Subgradient-based Subgradient Method, Convex Bundle Method,
Proximal Bundle Method

Gradient-based Gradient Descent, Conjugate Gradient, Stochastic,
Momentum, Nesterov, Averaged; Quasi-Newton with
(L-)BFGS, DFP, Broyden, SR1,...; Levenberg-Marquardt

Hessian-based Trust Regions, Adaptive Regularized Cubics (ARC)

splitting Chambolle-Pock, Douglas-Rachford, Cyclic Proximal Point,
Proximal Gradient

constrained Augmented Lagrangian, Exact Penalty, Frank-Wolfe,
Projected Gradient, Interior Point Newton

nonconvex Difference of Convex Algorithm, DCPPA

A Numerical Example



The Difference of Convex Algorithm in Manopt.jl

The algorithm is implemented and released in Julia using `Manopt.jl`¹. It can be used with any manifold from `Manifolds.jl`

A solver call looks like

```
q = difference_of_convex_algorithm(M, f, g, ∂h, p0)
```

where one has to implement `f(M, p)`, `g(M, p)`, and `∂h(M, p)`.

- ▶ a sub problem is generated if keyword `grad_g=` is set
- ▶ an efficient version of its cost and gradient is provided
- ▶ you can specify the sub-solver using `sub_state=`
to also set up the specific parameters of your favourite algorithm

Rosenbrock and First Order Methods

Problem. We consider the classical Rosenbrock example²

$$\arg \min_{x \in \mathbb{R}^2} a(x_1^2 - x_2)^2 + (x_1 - b)^2,$$

where $a, b > 0$, usually $b = 1$ and $a \gg b$, here: $a = 2 \cdot 10^5$.

Known Minimizer $x^* = \begin{pmatrix} b \\ b^2 \end{pmatrix}$ with cost $f(x^*) = 0$.

Goal. Compare first-order methods, e. g. using the (Euclidean) gradient

$$\nabla f(x) = \begin{pmatrix} 4a(x_1^2 - x_2) \\ -2a(x_1^2 - x_2) \end{pmatrix} + \begin{pmatrix} 2(x_1 - b) \\ 0 \end{pmatrix}$$



A “Rosenbrock-Metric” on \mathbb{R}^2

NTNU

In our Riemannian framework, we can introduce a new metric on \mathbb{R}^2 as

$$G_p := \begin{pmatrix} 1 + 4p_1^2 & -2p_1 \\ -2p_1 & 1 \end{pmatrix}, \text{ with inverse } G_p^{-1} = \begin{pmatrix} 1 & 2p_1 \\ 2p_1 & 1 + 4p_1^2 \end{pmatrix}.$$

We obtain $(X, Y)_p = X^\top G_p Y$

The exponential and logarithmic map are given as

$$\exp_p(X) = \begin{pmatrix} p_1 + X_1 \\ p_2 + X_2 + X_1^2 \end{pmatrix}, \quad \log_p(q) = \begin{pmatrix} q_1 - p_1 \\ q_2 - p_2 - (q_1 - p_1)^2 \end{pmatrix}.$$

`Manifolds.jl`:

Implement these functions on `MetricManifold(\mathbb{R}^2 , RosenbrockMetric())`.



The Riemannian Gradient w.r.t. the new Metric

NTNU

Let $f: \mathcal{M} \rightarrow \mathbb{R}$. Given the Euclidean gradient $\nabla f(p)$, its Riemannian gradient $\text{grad } f: \mathcal{M} \rightarrow T\mathcal{M}$ is given by

$$\text{grad } f(p) = G_p^{-1} \nabla f(p).$$

While we could implement this denoting $\nabla f(p) = (f'_1(p) \ f'_2(p))^\top$ using

$$\left\langle \text{grad } f(q), \log_q p \right\rangle_q = (p_1 - q_1)f'_1(q) + (p_2 - q_2 - (p_1 - q_1)^2)f'_2(q),$$

but it is automatically done in `Manopt.jl`.

The Experiment Setup

Algorithms. We now compare

1. The Euclidean gradient descent algorithm on \mathbb{R}^2 ,
2. The Riemannian gradient descent algorithm on \mathcal{M} ,
3. The Difference of Convex Algorithm on \mathbb{R}^2 ,
4. The Difference of Convex Algorithm on \mathcal{M} .

For DCA third we split f into $f(x) = g(x) - h(x)$ with

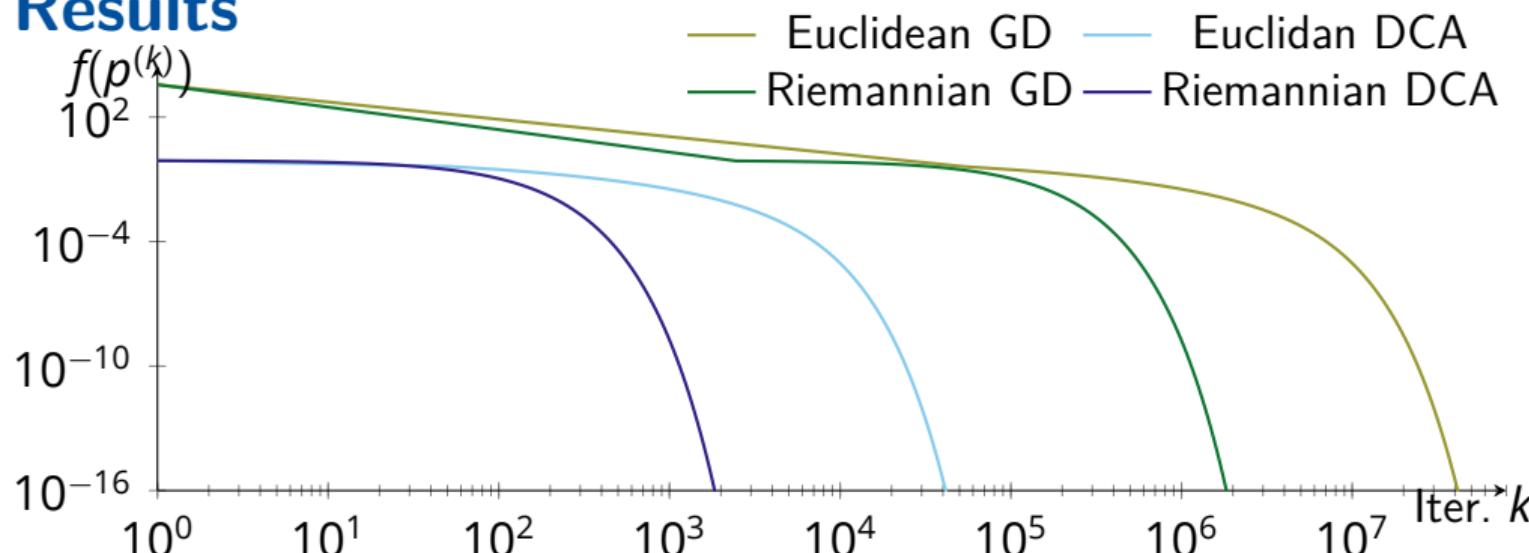
$$g(x) = a(x_1^2 - x_2)^2 + 2(x_1 - b)^2 \quad \text{and} \quad h(x) = (x_1 - b)^2.$$

Initial point. $p_0 = \frac{1}{10} \begin{pmatrix} 1 \\ 2 \end{pmatrix}$ with cost $f(p_0) \approx 7220.81$.

Stopping Criterion.

$$d_{\mathcal{M}}(p^{(k)}, p^{(k-1)}) < 10^{-16} \text{ or } \|\text{grad } f(p^{(k)})\|_p < 10^{-16}.$$

Results



Algorithm	Runtime (sec.)	# Iterations
Euclidean GD	305.567	53 073 227
Euclidean DCA	58.268	50 588
Riemannian GD	18.894	2 454 017
Riemannian DCA	7.704	2 459

Summary

Nonsmooth optimization on manifolds appears in several applications.

- ▶ many algorithms can be generalized
- ▶ many properties carry over, like convergence results
- ▶ Fenchel duality can be generalized [Schiela, Herzog, RB, 2024]
- ▶ Manifolds.jl & Manopt.jl [RB, 2022; Axen, Baran, RB, Rzecki, 2023]
- ▶ numerical examples available in ManoptExamples.jl

Next.

- ▶ LieGroups.jl [RB, Baran, 2026]
- ▶ ExponentialFamilyManifolds.jl and
ExponentialFamilyProjection.jl M. Lukashchuk, BIASlab, TU/e
based on ManifoldsBase.jl, using Manopt.jl

Selected References

-  Axen, S. D.; M. Baran; RB; K. Rzecki (2023). "Manifolds.jl: An Extensible Julia Framework for Data Analysis on Manifolds". *ACM Transactions on Mathematical Software* 49.4. DOI: [10.1145/3618296](https://doi.org/10.1145/3618296).
-  RB (2022). "Manopt.jl: Optimization on Manifolds in Julia". *Journal of Open Source Software* 7.70, p. 3866. DOI: [10.21105/joss.03866](https://doi.org/10.21105/joss.03866).
-  RB; M. Baran (2026). "Groups and smooth geometry using LieGroups.jl". *Proceedings of the JuliaCon Conferences*. Vol. 8. 79. *The Open Journal*, p. 195. DOI: [10.21105/jcon.00195](https://doi.org/10.21105/jcon.00195).
-  RB; O. P. Ferreira; E. M. Santos; J. C. d. O. Souza (2024). "The difference of convex algorithm on Hadamard manifolds". *Journal of Optimization Theory and Applications*. DOI: [10.1007/s10957-024-02392-8](https://doi.org/10.1007/s10957-024-02392-8).
-  RB; H. Jasa; P. John; M. Pfeffer (2025a). *The Intrinsic Riemannian Proximal Gradient Method for Nonconvex Optimization*. arXiv: [2506.09775](https://arxiv.org/abs/2506.09775).
-  — (2025b). *The Intrinsic Riemannian Proximal Gradient Method for Convex Optimization*. arXiv: [2507.16055](https://arxiv.org/abs/2507.16055).
-  RB; J. Persch; G. Steidl (2016). "A parallel Douglas Rachford algorithm for minimizing ROF-like functionals on images with values in symmetric Hadamard manifolds". *SIAM Journal on Imaging Sciences* 9.4, pp. 901–937. DOI: [10.1137/15M1052858](https://doi.org/10.1137/15M1052858).

